

pysib: An Open-Source Python Toolbox for SISO Polynomial System Identification

Diego Eckhard *

** Universidade Federal do Rio Grande do Sul, Porto Alegre, Brazil
(e-mail: diegoeck@ufrgs.br)*

Abstract: Discrete-time polynomial input-output models (ARX, ARMAX, OE, and Box-Jenkins) are usually estimated by prediction-error methods, but for OE, ARMAX, and BJ the finite-sample criterion is nonconvex: the estimate a user actually obtains is set by the initialization and the optimization procedure, not only by the asymptotic theory. This article documents the dedicated optimization strategy behind `pysib`, an open-source Python toolbox for SISO polynomial system identification—an ARX-based initialization, a smoothed-gradient phase, an incremental Gauss-Newton refinement, and filtered continuation interpreted as cost-function shaping. This strategy produced the filtered-continuation results reported by the author in earlier work but had not previously been described or released; it is given here in full and as open Python software, with a common five-polynomial representation, shared prediction and simulation routines, and the scripts and archived release needed to reproduce the experiments. On a moderate-noise OE benchmark the strategy returns estimates far more concentrated around the true parameters than a general-purpose nonlinear-programming solver, and on a harder nonconvex benchmark filtered continuation raises the success rate from 60% to 100%.

Keywords: system identification, Python, prediction-error method, polynomial models, open-source software

1. INTRODUCTION

Discrete-time polynomial input-output models remain a standard representation for linear system identification. In SISO applications, ARX, ARMAX, OE, and Box-Jenkins structures provide compact transfer-function descriptions in which the plant dynamics, input delay, and disturbance model can be represented by a small number of polynomials. This representation is attractive not only because of its long history in system identification, but also because it gives models that are directly usable for prediction, simulation, and control-oriented analysis (Ljung, 2010).

Prediction-error methods are commonly used to estimate these polynomial structures from input-output data. Under standard assumptions, including sufficient excitation, identifiability, stability, and model correctness, the global minimizer of the prediction-error criterion has the usual consistency and asymptotic covariance properties (Ljung, 1999; Söderström and Stoica, 1989). These properties, however, describe the ideal minimizer of the criterion. They do not by themselves specify how a finite-data, generally nonconvex optimization problem should be solved in software. For ARX this objective is linear least squares, but for OE, ARMAX, and BJ structures it is generally nonconvex because denominator and disturbance-model parameters enter the prediction recursion. Consequently, the numerical initialization and local optimization strategy become central to the estimator that the procedure yields in practice.

These numerical issues motivate software in which the model convention and the optimization procedure are

both explicit. A system-identification package should use a consistent representation across estimators and should make it possible to reproduce the numerical evidence supporting its use. `pysib` is an open-source Python toolbox for SISO polynomial system identification (Eckhard, 2026). Its scope is deliberately narrow: the package addresses discrete-time SISO polynomial input-output models, rather than MIMO identification, subspace methods, continuous-time identification, or nonlinear machine-learning models (Chen et al., 2012). Within this scope, it provides direct estimators, nonlinear PEM estimators, filtered variants, and a common prediction and simulation interface. This focus makes it possible to expose the polynomial structure used by the estimators while keeping the public interface compact.

Open Python tooling for this classical polynomial PEM setting is comparatively scarce, and the existing packages mostly address different problems. `SysIdentPy` targets NARMAX and other nonlinear model classes, with model structure selection and least-squares-type parameter estimation rather than the nonconvex transfer-function PEM considered here (Lacerda Junior et al., 2020). General-purpose control libraries such as `python-control` focus on analysis and design and offer only limited identification functionality (Fuller et al., 2021). The closest comparator is `SIPPY`, which does implement the ARX, ARMAX, OE, and BJ structures, but casts the polynomial criteria as a general nonlinear program solved with an off-the-shelf optimizer (Armenise et al., 2018). These tools are complementary rather than directly comparable; `pysib` occupies the specific niche of SISO polynomial PEM and pairs the

polynomial convention with an optimization procedure designed around the structure of the recursive PEM criterion, which is the part of the toolbox examined in most detail below.

The contribution of this article is threefold. First, it gives the first explicit description of the dedicated numerical strategy used by `pysib` for nonlinear PEM estimation in OE, ARMAX, and BJ structures. The strategy combines ARX-based initialization, a smoothed gradient phase, a Gauss–Newton refinement, and filtered continuation interpreted as cost-function shaping, in which the filters guide the optimizer through intermediate filtered criteria before returning to the original prediction-error problem. This strategy underlies the filtered-continuation results previously reported by the author (Eckhard et al., 2013, 2017), where it was used to produce comparisons against established identification toolboxes but was never itself documented. Second, the article provides an open-source Python implementation of the strategy, with a consistent polynomial model convention and a common prediction and simulation interface, addressing the scarcity of Python tooling for classical polynomial PEM. Third, it reports reproducible OE experiments that isolate the optimizer against a general-purpose solver and evaluate filtered continuation on a nonconvex benchmark. No new PEM theory is claimed; the emphasis is on disclosing and documenting an effective optimization strategy, releasing it as open software, and providing reproducible numerical evidence for the tested SISO OE settings.

The remainder of the article is organized as follows. Section 2 reviews the polynomial SISO model convention and PEM criterion. Section 3 describes the initialization, optimization, and filtered-continuation strategy. Section 4 summarizes the software interface and reproducibility information. Section 5 presents the numerical experiments, and the final section gives concluding remarks.

2. PREDICTION-ERROR IDENTIFICATION IN POLYNOMIAL SISO MODELS

Consider a discrete-time SISO data record $\{u(t), y(t)\}_{t=1}^N$. The objective is to estimate a finite-dimensional model of the relation between the measured input $u(t)$ and the measured output $y(t)$. It is useful to distinguish the deterministic response of the plant from the stochastic component of the measurement. The noise-free output is denoted by $y_0(t)$, and the measured output is written as the sum of this component and an additive disturbance:

$$y_0(t) = G_0(q^{-1})u(t), \quad (1)$$

$$y(t) = y_0(t) + v(t), \quad (2)$$

$$v(t) = H_0(q^{-1})e(t). \quad (3)$$

Here q^{-1} is the backward-shift operator, G_0 is the true plant, H_0 is the disturbance model, and $e(t)$ denotes the innovation process in the standard prediction-error setting. This notation separates the dynamics to be reproduced in simulation, represented by G_0 , from the dynamics used to describe the color of the disturbance, represented by H_0 .

The model family implemented in `pysib` follows the classical five-polynomial SISO convention (Ljung, 1999; Söderström and Stoica, 1989). This convention is useful because several standard structures can be expressed by

fixing selected polynomials to unity and estimating the remaining ones. With parameter vector θ and input delay n_k , the candidate model is represented as

$$y(t; \theta) = G(q^{-1}; \theta)u(t - n_k) + H(q^{-1}; \theta)e(t), \quad (4)$$

where

$$G(q^{-1}; \theta) = \frac{B(q^{-1}; \theta)}{A(q^{-1}; \theta)F(q^{-1}; \theta)}, \quad (5)$$

$$H(q^{-1}; \theta) = \frac{C(q^{-1}; \theta)}{A(q^{-1}; \theta)D(q^{-1}; \theta)}. \quad (6)$$

The polynomials are parameterized as

$$A(q^{-1}) = 1 + a_1q^{-1} + \dots + a_{n_a}q^{-n_a}, \quad (7)$$

$$B(q^{-1}) = b_1 + b_2q^{-1} + \dots + b_{n_b}q^{-(n_b-1)}, \quad (8)$$

$$C(q^{-1}) = 1 + c_1q^{-1} + \dots + c_{n_c}q^{-n_c}, \quad (9)$$

$$D(q^{-1}) = 1 + d_1q^{-1} + \dots + d_{n_d}q^{-n_d}, \quad (10)$$

$$F(q^{-1}) = 1 + f_1q^{-1} + \dots + f_{n_f}q^{-n_f}. \quad (11)$$

Thus, n_b denotes the number of numerator coefficients in B . In this conceptual parametrization, the input delay is not included in B ; it is the separate integer n_k appearing in (4). The transfer function $G(q^{-1}; \theta)$ represents the deterministic input–output dynamics, whereas $H(q^{-1}; \theta)$ represents the disturbance model used by the one-step-ahead predictor. In the model dictionary returned by the software, the same delay is stored by prepending n_k zeros to the polynomial B . This storage convention is used for simulation and filtering, but the order parameter n_b still counts only the nonzero numerator coefficients. In summary, the conceptual model carries the delay explicitly via n_k while B contains no delay; the software prepends n_k zeros to the returned B and does not expose n_k again through the model dictionary.

For a given θ , the one-step-ahead predictor associated with the model is the standard expression

$$\hat{y}(t | t-1; \theta) = H^{-1}(q^{-1}; \theta)G(q^{-1}; \theta)u(t - n_k) + [1 - H^{-1}(q^{-1}; \theta)]y(t), \quad (12)$$

which separates the deterministic filtered-input contribution from the autoregressive contribution of the past output. The prediction error and the finite-sample least-squares prediction-error criterion are

$$\varepsilon(t, \theta) = y(t) - \hat{y}(t | t-1; \theta), \quad (13)$$

$$V_N(\theta) = \frac{1}{N} \sum_{t=1}^N \varepsilon^2(t, \theta). \quad (14)$$

Thus, once a polynomial structure and model orders are selected, the estimation problem for the nonlinear structures is

$$\hat{\theta}_N = \arg \min_{\theta \in \Theta} V_N(\theta), \quad (15)$$

where Θ denotes the set of free coefficients of the chosen structure (for OE, $\theta = [b_1, \dots, b_{n_b}, f_1, \dots, f_{n_f}]$ with the constraint that the roots of the estimated polynomial F lie inside the unit circle). The nonlinear estimators in this article are numerical procedures for solving (15).

The motivation for using (14) is the usual asymptotic theory of prediction-error identification: under standard assumptions on identifiability, excitation, stability, and model correctness, the global minimizer has the classical consistency properties as $N \rightarrow \infty$, and a correctly parameterized disturbance model yields the standard asymp-

Table 1. Polynomial structures supported by pysib.

Structure	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>F</i>
ARX	free	free	1	1	1
ARMAX	free	free	free	1	1
OE	1	free	1	1	free
BJ	1	free	free	free	free

otic efficiency result (Ljung, 1999; Söderström and Stoica, 1989; Ljung, 1978). These statements concern the global minimizer and the large-sample limit; they do not imply that the criterion obtained from a finite data record is easy to minimize.

Table 1 summarizes the polynomial structures considered in the package. ARX fixes the noise model and is linear in the estimated parameters, so it can be computed by least squares and also serves as a natural direct estimate for initialization. OE fixes the disturbance model to white noise and estimates the plant numerator and denominator, which is appropriate when the main interest is the deterministic input–output dynamics. ARMAX adds a moving-average disturbance polynomial to the ARX structure, and BJ separates the plant and disturbance denominators, allowing a more flexible noise model.

These structural choices have direct numerical consequences. ARX leads to a linear least-squares problem. In contrast, OE, ARMAX, and BJ contain denominator and/or disturbance-model parameters inside the prediction recursion, so the resulting PEM criteria are generally nonconvex (Åström and Söderström, 1974; Söderström and Stoica, 1982). Their practical use therefore depends not only on writing down (14), but also on computing a useful local minimizer from finite data. The dedicated initialization and numerical optimization strategy used for this purpose is the subject of the next section.

3. NUMERICAL OPTIMIZATION AND INITIALIZATION STRATEGY

For the nonlinear structures in Section 2, the numerical procedure is an essential part of the practical estimator. The PEM criterion is defined by (14), but for OE, ARMAX, and BJ the finite-sample optimization problem is nonconvex and the result may depend on the initial point and on the path followed by the optimizer. A software implementation must therefore make concrete choices about initialization, descent safeguards, curvature information, and failure handling. The aim is not a new global-convergence guarantee, but a reliable and reproducible local-search strategy designed around the structure of the recursive PEM criterion for the SISO polynomial models considered by the toolbox (Goodwin et al., 2003).

The first design choice is to start each nonlinear search from an auxiliary ARX estimate. This gives a deterministic and inexpensive point in the parameter space before the recursive PEM criterion is optimized. Although the ARX model does not generally have the same disturbance structure as OE, ARMAX, or BJ, it provides a direct polynomial approximation with numerator and denominator dimensions compatible with the requested structure. The resulting coefficients initialize the corresponding plant

polynomials, while the remaining polynomial coefficients are set to neutral initial values. Thus, the nonlinear stage starts from a model that is dimensionally compatible with the requested structure, rather than from an arbitrary parameter vector.

Starting from this point, the optimizer proceeds in two phases. The reason for using two phases is that the information useful near a local minimizer is not necessarily reliable at the beginning of the search. A Gauss–Newton step is attractive when the current iterate is already in a region where the local least-squares approximation is informative, but it can be fragile when the initial estimate is still far from such a region or when the sensitivity matrix is ill conditioned. In preliminary testing, a direct damped Gauss–Newton or Levenberg–Marquardt iteration applied to the recursive PEM criterion was found unreliable under exactly these conditions, which motivated the two-phase strategy adopted here. The first phase is therefore deliberately conservative: it follows a smoothed gradient direction and uses small adaptive changes in the step length. This phase is intended to obtain a decrease of the criterion without first assuming that a local quadratic approximation is accurate. Let $g_k = \nabla V_N(\theta_k)$ denote the gradient of the criterion (14) evaluated at the current iterate. Instead of using g_k directly, the implementation updates a filtered direction

$$d_k = \frac{4d_{k-1} + g_k}{5}, \quad (16)$$

which damps abrupt changes induced by finite-data effects or by the recursive calculation of prediction errors. The smoothed direction is a conservative local search direction: it keeps memory of the previous direction while still responding to the current gradient. The trial point is then computed along the normalized direction,

$$\theta_{\text{trial}} = \theta_k - \alpha_k \frac{d_k}{\|d_k\|}. \quad (17)$$

The normalization separates the choice of direction from the scalar step length, so that the same step-length variable can be used even when the gradient magnitude changes substantially across iterations. If the trial point increases the cost, the step length is reduced; otherwise the point is accepted and the step length is increased:

$$\alpha_k \leftarrow \begin{cases} 0.99 \alpha_k, & \text{if } V_N(\theta_{\text{trial}}) > V_N(\theta_k), \\ 1.01 \alpha_k, & \text{otherwise.} \end{cases} \quad (18)$$

This rule is a numerical safeguard. It should be read as a practical descent heuristic, not as a global-convergence claim for the nonconvex PEM criterion.

After this initial decrease phase, the optimizer switches to a local refinement based on a Gauss–Newton approximation. At this stage, the algorithm uses the sensitivity of the prediction errors to parameter perturbations, which provides more directional information than the gradient alone when the current iterate is close enough to a meaningful basin of attraction. Let $S(\theta) \in \mathbb{R}^{N \times n_\theta}$ be the sensitivity matrix of the prediction errors with respect to the free parameters, with entries $S_{tj}(\theta) = \partial \varepsilon(t, \theta) / \partial \theta_j$. Up to scalar factors that do not affect the search direction, the Gauss–Newton curvature approximation is

$$H_{\text{GN}}(\theta) = S(\theta)^\top S(\theta), \quad (19)$$

and the search direction is obtained from

$$H_{\text{GN}}(\theta_k)p_k = g_k. \quad (20)$$

When $S(\theta_k)$ has full column rank, $H_{\text{GN}}(\theta_k)$ is positive definite. The system (20) is solved for p_k , and because it approximates $H_{\text{GN}}^{-1}g_k$, the Newton-like descent direction is $-p_k$; this is why the trial update shown next subtracts $\beta_i p_k$ rather than adding it. Because this approximation is most reliable only near a local minimizer, the implementation does not immediately take a full step. It instead forms trial points

$$\theta_{\text{trial}} = \theta_k - \beta_i p_k, \quad \beta_i \approx \frac{i}{1000}, \quad (21)$$

so that the effective Gauss–Newton step increases gradually over the iterations. This gradual transition is intended to combine the robustness of the previous descent phase with the faster local progress of curvature-based updates. In other words, the method does not abruptly replace the conservative phase by an aggressive full Newton step; it introduces curvature information progressively. If a trial point increases the cost or produces a nonfinite value, it is moved halfway back to the current iterate,

$$\theta_{\text{trial}} \leftarrow \frac{\theta_{\text{trial}} + \theta_k}{2}, \quad (22)$$

and the test is repeated a fixed number of times. The overall procedure is therefore a safeguarded local search: it first seeks a reliable decrease of the criterion and then uses sensitivity information to refine the estimate.

Each phase runs for a fixed iteration budget rather than a gradient-norm stopping test. The smoothed-gradient phase performs at most 10^4 step updates, organized as one hundred outer iterations of one hundred inner step-length adaptations, and exits early when the adaptive step length collapses below 10^{-7} . The Gauss–Newton phase then runs a fixed budget of one thousand iterations, over which the incremental factor β_i grows from near zero toward a full step. Fixing these budgets keeps the procedure deterministic and reproducible for a given data record.

The computationally intensive part of this process is the repeated evaluation of prediction errors, sensitivities, and gradients. These quantities are computed by sequential filtering recursions: each new sample depends on previous values of the filtered signals and sensitivity variables. Such recursions are not naturally expressible as a small number of vectorized array operations. For this reason, the inner recursions for OE, ARMAX, and BJ are implemented in C. This implementation choice keeps the recursion formulas close to the polynomial model structure and avoids Python loops in the numerical core, while the public estimator interface remains in Python.

The optimization procedure described so far acts within a single objective function: the two phases change how the search moves, but the criterion $V_N(\theta)$ defined by (14) stays the same throughout the call. In practice, however, the finite-sample shape of V_N can itself present the main difficulty, creating basins of attraction that lead a well-designed local search toward unacceptable local minima irrespectively of step-length choices. The filtered estimators described next address this problem at a different level: instead of modifying the optimizer, they modify the sequence of criteria presented to it.

3.1 Filtered Continuation as Cost-Function Shaping

The filtered estimators add a second mechanism for reducing the dependence on a single difficult finite-sample criterion. The preceding paragraphs describe how the optimizer moves within one criterion; filtered continuation acts at a different level by changing the sequence of criteria presented to the optimizer. It should not be interpreted merely as a preprocessing option applied before identification. If both input and output data are filtered and the PEM criterion is evaluated on the filtered record, the optimizer sees a different finite-sample objective. Each filter therefore defines an auxiliary criterion with its own landscape of local minima. Changing the filter changes the relative influence of different frequency regions and may make the early optimization problems easier than the original unfiltered problem.

This interpretation is the cost-function-shaping view of filtered OE identification (Eckhard and Bazanella, 2011; Eckhard et al., 2013, 2017). The purpose of the filters is to guide the optimizer through a sequence of related criteria in which undesirable local minima may have reduced influence. The wording is important: filtering can reshape the finite-sample criterion and can improve the numerical path followed by a local optimizer, but it does not remove the nonconvexity of the final PEM problem and does not provide a general guarantee of reaching its global minimum.

In `pysib`, the filtered variants implement this idea as a continuation procedure. A filtered identification problem is solved first; its solution is then used as the initial condition for the next, less restrictive filtered problem. The sequence continues until the final stage, which is the original unfiltered PEM criterion. Consequently, the returned estimate is still an estimate for the original model structure and criterion; the filters are used to shape the path toward that final problem. In the OE variant, nine filtered stages are applied before the final unfiltered OE call. Each stage uses a first-order low-pass filter whose pole is computed as

$$p = \exp(\log 0.05/(\tau \cdot 40)) \quad (23)$$

with $\tau = 0.9, 0.8, \dots, 0.1$; the resulting poles range from approximately 0.92 to 0.47, progressively relaxing the filtering constraint. In the ARMAX and BJ variants, nine filtered stages use first-order Butterworth filters with cutoff frequencies increasing from 0.1 to 0.9 of the Nyquist frequency. In all cases the continuation ends with the original unfiltered PEM criterion.

This mechanism is directly examined in the numerical experiments of Section 5. There, the standard OE estimator and the filtered OE continuation are applied to the same nonconvex benchmark, so that the effect of solving the intermediate filtered problems can be evaluated without changing the final simulation-error metric.

4. SOFTWARE INTERFACE AND REPRODUCIBILITY

The public interface of `pysib` is organized around the SISO input–output data record and the requested model orders. Each estimator returns a pair $(\hat{\theta}, m)$, where $\hat{\theta}$

Table 2. Public routines and their roles.

Routine	Structure or method	Role
arx	ARX	Linear/init.
sm	Stieglitz–McBride	OE init.
iv	Instrumental variables	ARX alternative
correlation	Correlation	ARX alternative
oe	OE	Nonlinear PEM
armax	ARMAX	Nonlinear PEM
bj	Box–Jenkins	Nonlinear PEM
*filtered	Filtered continuation	Cost shaping
predict/simulate	Model evaluation	Common evaluation

is the vector of free estimated parameters and m is a structured polynomial representation of the identified model. This convention separates the numerical parameter vector used by the optimizer from the model object used for prediction, simulation, and comparison. The vector $\hat{\theta}$ is useful when the user wants to inspect the numerical parameters directly, for example in Monte Carlo studies of parameter dispersion. The structured object m , in contrast, is the representation used by the rest of the package to evaluate the identified model.

The model representation m is a dictionary with keys **A**, **B**, **C**, **D**, and **F**. These entries correspond directly to the five polynomials introduced in Section 2; polynomials that are fixed by a particular structure are represented by the scalar polynomial [1]. The same convention is used by the linear estimators, the nonlinear PEM estimators, and the filtered variants. As a result, changing the estimator does not require changing the downstream code used for prediction, simulation, or error computation.

Table 2 summarizes the main routines. The first group contains direct or linear methods that can be used either as estimators in their own right or as auxiliary initial conditions. The **sm** routine implements the Stieglitz–McBride method (Stieglitz and McBride, 1965). The nonlinear routines **oe**, **armax**, and **bj** call the PEM optimizer described in Section 3. The filtered variants expose the continuation strategy of Section 3.1 through the same model convention. This common interface is also important for the experiments in Section 5, because all methods can be evaluated by the same simulation routine and the same error metric.

The direct methods **iv**, **correlation**, and **sm** do not rely on a PEM optimizer and are included as diagnostic and initialization tools. Instrumental variables reduce the bias introduced by colored output noise by using auxiliary instruments, typically at the cost of increased variance compared to least squares. The correlation estimator provides an alternative ARX estimate based on sample correlations rather than the normal equations. The Stieglitz–McBride method (Stieglitz and McBride, 1965) is a classical iterative procedure for OE models that can diverge in some configurations but is often effective as a fast initializer for nonlinear PEM. All three estimators can be used as stand-alone routines and their output models can also supply initial parameter vectors for the **oe**, **armax**, and **bj** routines.

The following minimal example illustrates the common interface for OE identification and model evaluation:

```
import pysib

theta, m = pysib.oe(u, y, nb=1, nf=3, nz=1)
```

```
yp = pysib.predict(u, y, m)
ys = pysib.simulate(u, m)
```

Here u and y are the measured input and output vectors. The call to **oe** estimates an OE model with one numerator coefficient, a third-order F polynomial, and one sample of input delay. The same returned model representation m is then used for one-step-ahead prediction and for noise-free simulation. The same pattern applies to models returned by the other estimators: once m has been constructed, the evaluation code is independent of the routine that produced it.

Any model returned by an estimator can be passed to the common **predict** and **simulate** routines. The former computes one-step-ahead predictions using the model disturbance structure, whereas the latter computes the simulated noise-free response of the plant model. Keeping these operations separate is important in system-identification experiments, since prediction and simulation measure different aspects of the identified model. This evaluation layer is also the mechanism used to compute the simulation-error metric in the numerical experiments of Section 5.

Formally, given a model m returned by any estimator, the **predict** routine evaluates the one-step-ahead predictor (12) using the polynomials in m , with G and H formed from those polynomials and the input delay included in B according to the storage convention. The noise-free simulation computed by **simulate** uses only the plant model,

$$\hat{y}_s(t) = G(q^{-1})u(t). \quad (24)$$

Thus, **predict** uses the full model (plant and disturbance), while **simulate** uses only the deterministic plant dynamics.

Table 3. Software availability.

Version	pysib v0.2.1 (Eckhard, 2026)
License	MIT
Install	<code>pip install pysib</code>
Documentation	https://pysib.net/
Source	https://github.com/diegoeck/pysib
DOI	https://doi.org/10.5281/zenodo.20572074
Tests and experiments	Included in repository

Detailed API documentation is left to the accompanying software documentation; the purpose of this article is to describe the numerical methods and the evidence obtained from the reproducible experiments.

5. NUMERICAL EXPERIMENTS

The numerical experiments are restricted to SISO OE examples. This choice keeps the empirical analysis focused on plant-dynamics estimation and on the numerical behavior of the PEM optimizer, without adding the separate issue of estimating a colored-noise model. OE is also a direct representative of the nonconvex polynomial PEM problems targeted by the dedicated optimization strategy. The two experiments are designed to probe that strategy: the first contrasts the dedicated PEM optimizer with a general-purpose solver on a benign problem, and the second isolates the effect of filtered continuation on a harder

nonconvex problem. Together they provide empirical evidence for the design choices of Section 3 in the tested OE settings. Broader comparisons of the filtered method against established MATLAB identification toolboxes are reported in earlier work (Eckhard et al., 2013, 2017); the present experiments instead isolate the optimizer and the continuation within a single reproducible toolbox. The ARMAX and BJ structures use the same optimization strategy and interface; preliminary experiments not reported here suggested that it is effective for them as well, and further examples are straightforward to produce within the same framework but are omitted here for space.

Both experiments use the same third-order OE plant. The conceptual parametrization of Section 2 gives

$$\begin{aligned} B(q^{-1}) &= 1, \quad n_k = 1, \\ F(q^{-1}) &= 1 - 2.4q^{-1} + 1.91q^{-2} - 0.504q^{-3}. \end{aligned} \quad (25)$$

The stored polynomial representation used by the software (in which the unit delay is encoded by a leading zero) is

$$\begin{aligned} B(q^{-1}) &= q^{-1}, \\ F(q^{-1}) &= 1 - 2.4q^{-1} + 1.91q^{-2} - 0.504q^{-3}. \end{aligned} \quad (26)$$

This stored form is the one used to simulate the plant and to compute the simulation error. For each experiment, the noise-free output y_0 is generated by simulating (26) with the specified input sequence, and the measured output is

$$y(t) = y_0(t) + v(t), \quad (27)$$

$$v(t) \sim \mathcal{N}(0, \sigma_v^2). \quad (28)$$

The sample index is $t = 1, \dots, N$, and the random seed is fixed to zero in the experiment scripts. The main metric is the relative simulation error

$$E_{\text{sim}} = \frac{\|y_0 - \hat{y}\|_2}{\|y_0\|_2}. \quad (29)$$

Here y_0 is the noise-free plant output and $\hat{y}$ is the simulated output of the estimated model. This metric evaluates the estimated plant dynamics rather than the one-step-ahead predictor, which is appropriate for the OE benchmarks considered here. A run is counted as successful when $E_{\text{sim}} < 5\%$.

5.1 Noisy-Data Comparison with SIPPY

The first experiment compares `pysib` with the SIPPY OE implementation (Armenise et al., 2018) (the `sippy_unipi` distribution, version 1.0.1, called through its `system_identifi` routine with the OE method and matching model orders), using noisy data generated from (26). The two implementations take different routes to the same OE structure: SIPPY formulates the problem as a general nonlinear program and solves it with an interior-point optimizer, whereas `pysib` applies the dedicated PEM strategy of Section 3 from an ARX initialization. The comparison is therefore between a general-purpose solver and an optimizer built around the recursive PEM criterion. The noise level $\sigma_v = 1$ defines a moderate-noise setting in which all methods are expected to identify the plant reasonably well, so the relevant quantities are the concentration of the estimated parameters and the distribution of the simulation error. In other words, this benchmark does not test whether the plant is identifiable, but how tightly the underlying optimization converges to the data-generating parameters in a benign, near-convex case. The Monte

Carlo study uses $M = 100$ independent data records of length $N = 1000$. Additive white output noise is added to the noise-free output. The input is the deterministic multisine

$$u(t) = \sin(2\pi t/17) + 0.7 \sin(2\pi t/43) + 0.4 \sin(2\pi t/97). \quad (30)$$

The methods compared are the Stieglitz–McBride estimator in `pysib`, the OE PEM estimator in `pysib`, and the SIPPY OE estimator, all using the same OE order. The parameter table is included to show whether the estimates are centered near the coefficients of the data-generating plant, while the error table and figure show the effect of these estimates on simulated output.

Table 4. Parameter statistics for the noisy-data experiment.

Method	Parameter	Mean	Std
SM	b_1	1.0005	0.0030
SM	f_1	-2.3997	0.0018
SM	f_2	1.9094	0.0033
SM	f_3	-0.5037	0.0015
OE	b_1	1.0005	0.0030
OE	f_1	-2.3997	0.0018
OE	f_2	1.9094	0.0033
OE	f_3	-0.5037	0.0015
SIPPY	b_1	1.0439	0.0586
SIPPY	f_1	-2.3731	0.0342
SIPPY	f_2	1.8604	0.0630
SIPPY	f_3	-0.4810	0.0292

Table 5. Simulation-error statistics for the noisy-data experiment.

Method	Mean error	Std error	Success
SM	1.05×10^{-3}	3.88×10^{-4}	100%
OE	1.05×10^{-3}	3.89×10^{-4}	100%
SIPPY	9.71×10^{-3}	8.97×10^{-3}	100%




Fig. 1. Distribution of the relative simulation error for `pysib` SM, `pysib` OE, and SIPPY OE in the noisy-data comparison.

Table 4 reports the empirical means and standard deviations of the estimated OE parameters, while Table 5 and Figs. 1 and 2 summarize the simulation errors. Because this is a benign, near-convex configuration, the agreement between the two `pysib` estimators is itself informative: the Stieglitz–McBride estimator and the ARX-initialized OE PEM estimator coincide to the reported precision and sit




Fig. 2. Boxplot of the relative simulation error for the same methods.

essentially at the data-generating parameters, indicating that two independent search paths reach the same minimizer. The SIPPY OE estimates obtained in this configuration, by contrast, exhibit a visible bias—for example, a mean b_1 of 1.044 against the true value 1.0—together with a standard deviation roughly an order of magnitude larger and a mean simulation error about nine times larger. The histogram and boxplot show the same effect from the model-output perspective: the `pysib` error distributions are tighter and closer to zero. All methods satisfy the success criterion in all Monte Carlo runs, so the distinction in this moderate-noise case is not whether the plant can be identified, but how tightly the underlying optimization converges to the true minimizer.

5.2 Filtered Continuation on a Nonconvex OE Problem

The second experiment evaluates the filtered-continuation mechanism described in Section 3.1. The plant is again (26), with $M = 100$ Monte Carlo runs and data length $N = 1000$, but the additive white output noise standard deviation is increased to $\sigma_v = 30$. This much larger noise level is used to create a more challenging nonconvex optimization problem, in which the standard OE search is more likely to converge to an undesirable local minimum. The input is

$$u(t) = \sin(2\pi t/18) + \sin(2\pi t/28) + \sin(2\pi t/61). \quad (31)$$

The comparison is between the standard `pysib` OE estimator and the filtered OE continuation. Both methods are evaluated by the same final simulation-error metric (29); this is important because the filters are used only to guide the optimization path, not to change the final model-evaluation criterion.

The results in Table 6 and Figs. 3 and 4 show a clear empirical difference between the two procedures in this benchmark. Standard OE attains the success criterion in

Table 6. Simulation-error statistics for the filtered-continuation experiment.

Method	Mean error	Std error	Success
OE	5.69×10^{-2}	4.37×10^{-2}	60%
OE filtered	2.17×10^{-2}	7.94×10^{-3}	100%




Fig. 3. Distribution of the relative simulation error for standard OE and filtered OE in the nonconvex benchmark.




Fig. 4. Boxplot of the relative simulation error for the same methods.

60% of the runs, whereas the filtered continuation attains it in all runs and also reduces the mean relative simulation error. The histogram illustrates that the improvement is not only an average effect: the filtered continuation substantially reduces the high-error tail observed for standard OE in this experiment. This supports the role of filtered continuation as a cost-shaping mechanism for this OE problem. The conclusion is deliberately empirical: the experiment demonstrates improved behavior in the tested

SISO OE setting, but it does not imply a global optimality guarantee for nonconvex PEM.

6. CONCLUSION

This article presented `pysib`, an open-source Python toolbox for SISO polynomial system identification. Three aspects were described: a dedicated numerical strategy for nonlinear PEM estimation, the software implementation that exposes this strategy through a compact public interface, and the numerical experiments used to examine the behavior of the toolbox.

The numerical strategy combines several complementary elements. An auxiliary ARX estimate supplies a deterministic initial point with dimensions compatible with the requested nonlinear structure. A smoothed-gradient phase then provides a conservative initial decrease of the PEM criterion, before an incremental Gauss–Newton phase introduces curvature information for local refinement. The filtered variants add a separate continuation mechanism: instead of changing only the step taken by the optimizer, they change the sequence of finite-sample criteria presented to it and then return to the original unfiltered PEM problem.

The software contribution is to expose this strategy through a compact Python interface while keeping a consistent polynomial model representation across estimators. The returned pair $(\hat{\theta}, m)$ separates parameter inspection from model evaluation, and the same `predict` and `simulate` routines can be applied to models produced by different estimators. Together with the PyPI release, public repository, archived DOI, experiment scripts, and automated tests, this interface is intended to make the numerical results reproducible rather than dependent on undocumented implementation details.

The numerical evidence focused on OE benchmarks. In the noisy-data comparison, the `pysib` SM and OE estimates were tightly concentrated around the true parameters and attained a small mean simulation error. In the nonconvex OE experiment, filtered continuation improved the empirical success rate from 60% for standard OE to 100% for filtered OE, supporting its role as a cost-shaping mechanism in that setting. These results should be read as reproducible evidence for the tested SISO OE problems, not as a general performance claim for all polynomial identification settings.

The scope is deliberately limited. The toolbox and experiments considered here address SISO linear polynomial input–output models; they do not cover MIMO identification, subspace methods, continuous-time models, or nonlinear machine-learning models. Moreover, the optimizer is a practical local method for nonconvex PEM criteria and does not provide a global-optimality guarantee.

FUTURE WORK

Several directions follow naturally from the current scope. The most immediate extension is to MIMO polynomial input–output models, where the five-polynomial convention and the common prediction and simulation interface could be generalized to multi-variable structures. A systematic experimental evaluation of the ARMAX and BJ

routines—already implemented with the same optimization strategy described in Section 3—is also needed to complement the OE benchmarks reported here. Further work on filtered continuation could explore adaptive filter schedules or filter families beyond first-order low-pass and Butterworth designs. Finally, broader comparisons with other Python identification packages and with the MATLAB System Identification Toolbox would help position `pysib` in the wider ecosystem.

DECLARATION OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES IN THE MANUSCRIPT PREPARATION PROCESS

During the preparation of this work the author used Claude (Anthropic), GPT (OpenAI), and DeepSeek in order to improve the clarity and language of the manuscript and to assist with software development. After using these tools/services, the author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

REFERENCES

- Armenise, G., Vaccari, M., Di Capaci, R.B., and Pannocchia, G. (2018). An open-source system identification package for multivariable processes. In *2018 UKACC 12th International Conference on Control (CONTROL)*. doi:10.1109/CONTROL.2018.8516791. URL <https://doi.org/10.1109/CONTROL.2018.8516791>.
- Åström, K.J. and Söderström, T. (1974). Uniqueness of the maximum likelihood estimates of the parameters of an ARMA model. *IEEE Transactions on Automatic Control*, 18(6), 769–773. doi:10.1109/TAC.1974.1100795.
- Chen, T., Ohlsson, H., and Ljung, L. (2012). On the estimation of transfer functions, regularizations and Gaussian processes—revisited. *Automatica*, 48(8), 1525–1535. doi:10.1016/j.automatica.2012.05.026.
- Eckhard, D. and Bazanella, A.S. (2011). On the global convergence of identification of output error models. In *Proc. 18th IFAC World Congress*. Milan, Italy. doi:10.3182/20110828-6-IT-1002.03566.
- Eckhard, D., Bazanella, A.S., Rojas, C.R., and Hjalmarsson, H. (2013). Input design as a tool to improve the convergence of pem. *Automatica*, 49(11), 3282–3291. doi:10.1016/j.automatica.2013.08.027. URL <https://doi.org/10.1016/j.automatica.2013.08.027>.
- Eckhard, D., Bazanella, A.S., Rojas, C.R., and Hjalmarsson, H. (2017). Cost function shaping of the output error criterion. *Automatica*, 76, 53–60. doi:10.1016/j.automatica.2016.10.015. URL <https://doi.org/10.1016/j.automatica.2016.10.015>.
- Eckhard, D. (2026). `pysib`: System identification toolbox. Version 0.2.1. doi:10.5281/zenodo.20572074. URL <https://doi.org/10.5281/zenodo.20572074>.
- Fuller, S., Greiner, B., Moore, J., Murray, R.M., van Paassen, R., and Yorke, R. (2021). The Python control systems library (python-control). In *2021 60th IEEE Conference on Decision and Control (CDC)*, 4875–4881. doi:10.1109/CDC45484.2021.9683368.
- Goodwin, G.C., Agüero, J.C., and Skelton, R. (2003). Conditions for local convergence of maximum likelihood estimation for ARMAX models. In *Proc. 13th IFAC*

- Symposium on System Identification*, 797–802. Rotterdam, The Netherlands.
- Lacerda Junior, W.R., Martins, S.A.M., and Nepomuceno, E.G. (2020). SysIdentPy: A Python package for system identification using NARMAX models. *Journal of Open Source Software*, 5(54), 2384. doi:10.21105/joss.02384.
- Ljung, L. (1978). Convergence analysis of parametric identification methods. *IEEE Transactions on Automatic Control*, 23(5), 770–783. doi:10.1109/TAC.1978.1101840.
- Ljung, L. (2010). Perspectives on system identification. *Annual Reviews in Control*, 34(1), 1–12. doi:10.1016/j.arcontrol.2009.12.001.
- Ljung, L. (1999). *System Identification: Theory for the User*. Prentice Hall, Upper Saddle River, NJ, 2 edition.
- Söderström, T. and Stoica, P. (1982). Some properties of the output error method. *Automatica*, 18(1), 93–99. doi:10.1016/0005-1098(82)90031-0.
- Söderström, T. and Stoica, P. (1989). *System Identification*. Prentice Hall, Englewood Cliffs, NJ.
- Stieglitz, K. and McBride, L.E. (1965). A technique for the identification of linear systems. *IEEE Transactions on Automatic Control*, 10(4), 461–464. doi:10.1109/TAC.1965.1098181. URL <https://doi.org/10.1109/TAC.1965.1098181>.